

TOP 100

AZURE

DATA

FACTORY

INTERVIEW

PREPARATION

QUESTIONS AND

ANSWERS

3+ YOE

BASED ON REAL-TIME

SCENARIOS, 2025

Section 1 – Core Azure Data Factory (ADF) Scenarios(1-30)

Scenario 1 – Incremental Data Load Using Watermarking

Description:

Load only new or updated records from a SQL source to Azure Data Lake.

Explanation:

Use a high-watermark column such as `LastModifiedDate`. Store the latest value in a control table or file and use it in your next run.

Example:

```
SELECT * FROM Sales  
WHERE LastModifiedDate > '@{pipeline().parameters.LastWatermark}'
```

Flow: Source → Copy Activity → ADLS → Update Watermark

Key Benefit: Loads delta data only, improving performance.

Scenario 2 – Parameterizing Pipelines for Reusability

Description:

Make one pipeline handle multiple tables dynamically.

Explanation:

Add dataset parameters (`TableName`, `SinkPath`) and pass them through pipeline parameters using dynamic expressions.

Example: `@concat('/raw/', pipeline().parameters.TableName)`

Flow: Lookup → ForEach → Copy Activity

Key Benefit: Reduces duplication and simplifies maintenance.

Scenario 3 – Failure Alerts via Logic App

Description:

Send an email if any activity fails.

Explanation:

Use an “On Failure” dependency with a Web Activity calling a Logic App endpoint.

Flow: ADF → Web Activity → Logic App → Email

Key Benefit: Immediate awareness of pipeline failures.

Scenario 4 – Event and Schedule Triggers

Description:

Run daily loads at 2 AM or when new files arrive.

Explanation:

Configure both Schedule and Blob Event triggers pointing to the same pipeline.

Key Benefit: Hands-free automation and timely refreshes.

Scenario 5 – Dynamic Table Loads Using Lookup and ForEach

Description:

Load multiple tables defined in a config file.

Explanation:

Read the config via Lookup, then ForEach loop over each record, executing Copy Activity with parameters.

Flow: Lookup → ForEach → Copy

Key Benefit: Scales easily for hundreds of tables.

Scenario 6 – Error Handling with Custom Logging

Description:

Capture activity errors in a custom log table.

Explanation:

Use “On Failure” paths to execute a Stored Procedure Activity inserting pipeline name, activity name, and error message into SQL log table.

Flow: Activity → OnFailure → Stored Procedure

Key Benefit: Centralized error tracking for monitoring.

Scenario 7 – Copying Data Between Different Regions

Description:

Move data from East US SQL DB to West Europe ADLS.

Explanation:

Use two self-hosted IRs and enable staging via Blob Storage to optimize cross-region transfer.

Flow: Source → Staging Blob → Sink

Key Benefit: Faster and secure inter-region transfers.

Scenario 8 – Using Variables and Set Variable Activity

Description:

Maintain loop counters or dynamic file paths within a pipeline.

Explanation:

Initialize a variable and update its value inside loops using Set Variable Activity.

Example: @concat('/archive/', item().FileName)

Key Benefit: Better control flow and flexibility.

Scenario 9 – Validating Input Files Before Processing

Description:

Ensure files follow correct naming and schema before copying.

Explanation:

Use Get Metadata Activity to check file size and name patterns, then If Condition Activity to proceed or fail.

Flow: Get Metadata → If Condition → Copy

Key Benefit: Prevents bad data loads.

Scenario 10 – Data Flow for Complex Transformations

Description:

Perform joins, aggregations, and derived columns without Databricks.

Explanation:

Use ADF Mapping Data Flows with Source, Join, Aggregate, and Sink transformations.

Flow: ADLS → Data Flow → Synapse Table

Key Benefit: Low-code ETL within ADF for complex logic.

💡 Mini Interview Tips (for Scenarios 1–10)

- Stress the use of dynamic content and parameters.
 - Mention how you track incremental loads and errors.
 - Emphasize monitoring and automation via triggers and Logic Apps.
-

Scenario 11 – Data Archival After Successful Load

Description:

After ingesting data from the source, archive the original file to a backup folder.

Explanation:

Use a Copy Activity to move the file, followed by a Delete Activity to remove it from the incoming folder.

Configure “On Success” dependency from Copy → Delete.

Flow: ADLS /Input → Copy → ADLS /Archive

Key Benefit: Keeps source folders clean and prevents duplicate loads.

Scenario 12 – Using Lookup for Dynamic SQL

Description:

Store complex SQL queries in a config table and execute them dynamically.

Explanation:

The Lookup Activity fetches SQL statements from control DB. Use

`@activity('Lookup1').output.firstRow.Query` inside Copy Activity.

Flow: Lookup (Config Table) → Copy Activity

Key Benefit: Externalizes logic for easier maintenance.

Scenario 13 – Parallel Copy Execution

Description:

Run multiple Copy Activities in parallel for faster ingestion.

Explanation:

Use ForEach Activity with “Batch count” set > 1 to execute concurrent copies.

Flow: ForEach (Parallel) → Copy Activities

Key Benefit: Improves performance and reduces overall pipeline runtime.

Scenario 14 – Implementing Retry Logic

Description:

Handle transient connectivity issues gracefully.

Explanation:

Set the “Retry” property in activities (e.g., 3 attempts, 30 seconds interval).

Key Benefit: Increases reliability for unstable network or service connections.

Scenario 15 – Using Stored Procedure for Post-Load Validation

Description:

Validate target table row count after loading.

Explanation:

After Copy Activity, call a Stored Procedure Activity that checks `source count = target count`.

Flow: Copy → Stored Procedure (Validation)

Key Benefit: Ensures data integrity before marking job complete.

Scenario 16 – Filter and Conditional Execution

Description:

Run only specific files that meet a condition.

Explanation:

Use “Filter Activity” on metadata to select items where, e.g., `FileSize > 0`. Pass result to ForEach.

Flow: Get Metadata → Filter → ForEach → Copy

Key Benefit: Avoids unnecessary processing of empty files.

Scenario 17 – Dynamic Folder Creation in ADLS

Description:

Create a new folder structure each day automatically.

Explanation:

Use `@concat('/raw/', pipeline().parameters.LoadDate)` in Sink dataset.

Key Benefit: Organizes data by date for easier tracking and purging.

Scenario 18 – Pipeline Failure Dependency Handling

Description:

Stop downstream pipelines if upstream fails.

Explanation:

Use “On Failure” dependencies or trigger pipelines with run-status checks through REST API.

Key Benefit: Prevents cascading errors and maintains data consistency.

Scenario 19 – Custom Logging to Azure Monitor

Description:

Push pipeline run status and duration to Azure Log Analytics.

Explanation:

Call Azure Monitor Data Collector API using Web Activity with pipeline run details.

Key Benefit: Centralized monitoring and trend analysis for ADF runs.

Scenario 20 – Data Validation Before Insert

Description:

Reject bad records before inserting into target.

Explanation:

Use Mapping Data Flow with Derived Column and Conditional Split transformations to separate valid and invalid data.

Flow: Source → Conditional Split → Sink (Valid/Invalid)

Key Benefit: Ensures clean data in target systems.

💡 Mini Interview Tips (Scenarios 11 – 20)

- Mention how you handle parallelism, error retries, and data validation.
 - Stress real-time logging and monitoring using Azure Monitor.
 - Explain how archival and folder automation help with data governance.
-

Scenario 21 – Incremental Data Load Using Watermark

Description:

Load only new or updated data instead of full load.

Explanation:

Maintain a control table with last load timestamp. Use it in the source query:

```
SELECT * FROM Sales WHERE ModifiedDate > @{variables('LastWatermark')}
```

Update watermark after load.

Key Benefit: Reduces data movement and improves efficiency.

Scenario 22 – Parameterized Linked Services

Description:

Use a single Linked Service for multiple databases or environments.

Explanation:

Enable “Parameterize” in Linked Service and pass values like `ServerName`, `DBName` dynamically.

Key Benefit: Simplifies deployment and avoids duplicate configurations.

Scenario 23 – Multiple File Merge Before Load

Description:

Combine multiple CSVs into one before ingestion.

Explanation:

Use Data Flow’s *Union* transformation or use Copy Activity with wildcard path (e.g., `input/*.csv`).

Key Benefit: Reduces number of small files and simplifies load management.

Scenario 24 – File Arrival Trigger

Description:

Automatically trigger pipeline when a file lands in storage.

Explanation:

Use *Event-based trigger* with Azure Blob Storage event grid integration.

Key Benefit: Enables near real-time ingestion without manual intervention.

Scenario 25 – Handling CSV with Variable Columns

Description:

Source CSVs may have different schema structures.

Explanation:

Enable “Skip line count = 1” and use *Schema Drift* in Data Flow. Add Derived Columns to standardize schema.

Key Benefit: Supports semi-structured ingestion flexibly.

Scenario 26 – Securely Storing Credentials

Description:

Avoid hardcoding connection strings.

Explanation:

Integrate Azure Key Vault with Linked Services. Use `@Microsoft.KeyVault(SecretName)` reference.

Key Benefit: Enhances security and simplifies credential rotation.

Scenario 27 – Dynamic Pipeline Execution from Master Pipeline

Description:

Execute child pipelines dynamically based on configuration.

Explanation:

Use Lookup to fetch pipeline names from config table, ForEach to iterate, and *Execute Pipeline Activity* to trigger.

Key Benefit: Enables metadata-driven orchestration.

Scenario 28 – Sending Pipeline Alerts via Email

Description:

Notify team on failure or success.

Explanation:

Use Logic App or Web Activity to send email notification with run status.

Key Benefit: Improves visibility and proactive issue resolution.

Scenario 29 – Data Transformation with Databricks

Description:

Perform heavy data cleansing and joins using Databricks Notebook.

Explanation:

Use *Databricks Notebook Activity* in ADF pipeline and pass parameters dynamically.

Example:

```
df = spark.read.csv(input_path)
df_clean = df.filter(df["amount"] > 0)
df_clean.write.parquet(output_path)
```

Key Benefit: Handles complex transformations at scale.

Scenario 30 – Copying Data Between Subscriptions

Description:

Migrate data from one Azure subscription to another.

Explanation:

Use *Managed Identity* authentication and configure *Cross-Subscription Linked Services*.

Flow: ADF → Source Blob (Subscription A) → Target Blob (Subscription B)

Key Benefit: Simplifies enterprise-level data movement.

💡 Mini Interview Tips (Scenarios 21 – 30)

- Focus on **incremental load** and **metadata-driven pipelines**, as these are real-world must-haves.
 - Emphasize **Key Vault integration** for secure, enterprise-grade design.
 - Be ready to explain **Databricks collaboration** for advanced ETL.
-

Section 2 – ADF + Databricks Integration Scenarios (31 to 50).

Scenario 31 – ADF Triggering Databricks Notebook

Description:

ADF calls a Databricks notebook for transformation.

Explanation:

Use *Databricks Notebook Activity* in ADF with parameters for path, cluster, and notebook name.

Example:

ADF → Notebook → Reads data from ADLS → Writes cleaned parquet back.

Key Benefit: Combines orchestration (ADF) + transformation (Databricks) in one flow.

Scenario 32 – Parameter Passing Between ADF and Databricks

Explanation:

Define parameters in Databricks notebook, e.g.:

```
dbutils.widgets.text("SourcePath", "")  
src=dbutils.widgets.get("SourcePath")
```

ADF passes dynamic values through activity settings.

Key Benefit: Promotes reusable notebooks with environment-based inputs.

Scenario 33 – Mounting ADLS in Databricks

Explanation:

Mount ADLS storage once for easy read/write:

```
dbutils.fs.mount(  
  "wasbs://container@storage.blob.core.windows.net/",  
  "/mnt/data")
```

Benefit: Simplifies data access for multiple notebooks.

Scenario 34 – Incremental Load with Databricks Merge

Explanation:

Use Delta Lake's MERGE to update target incrementally:

```
MERGE INTO target USING source  
ON target.id=source.id  
WHEN MATCHED THEN UPDATE SET *  
WHEN NOT MATCHED THEN INSERT *
```

Benefit: Fast upserts vs full refresh.

Scenario 35 – ADF Pipeline to Run Databricks Job Cluster

Explanation:

ADF notebook activity can specify job cluster config (JSON).

Automatically creates and deletes cluster post-execution.

Benefit: Cost control and auto-scaling for ETL.

Scenario 36 – Error Handling Between ADF and Databricks

Explanation:

Use `try-except` in Databricks and throw custom errors:

```
try:  
    df=spark.read.csv(path)  
except Exception as e:  
    dbutils.notebook.exit(str(e))
```

ADF captures failure in *On Failure* path.

Benefit: Improves debugging and logging.

Scenario 37 – Data Quality Check in Databricks

Explanation:

Notebook validates row counts and nulls before load.

Example:

```
bad=df.filter(df["amount"].isNull())  
if bad.count()>0:  
    raise Exception("Null values found")
```

Benefit: Prevents bad data from propagating.

Scenario 38 – ADF Lookup to Fetch Dynamic Notebook Path

Explanation:

Store notebook paths in SQL config table.

Lookup activity + ForEach calls Databricks Notebook dynamically.

Benefit: Metadata-driven automation.

Scenario 39 – Integrating Databricks Delta with Power BI

Explanation:

After ADF→Databricks pipeline creates Delta tables, connect Power BI directly to them using *Azure Synapse Serverless* or *ODBC*.

Benefit: Near real-time analytics on refreshed data.

Scenario 40 – Partitioned Parquet Output

Explanation:

Databricks writes partitioned data:

```
df.write.partitionBy("year", "month").parquet(output)
```

ADF reads from dynamic folder path.

Benefit: Optimizes read performance and query pushdown.

Scenario 41 – Automating Cluster Start/Stop via ADF

Explanation:

Use Databricks REST API in ADF Web activity to start/stop clusters before/after jobs.

Benefit: Cost optimization for ad-hoc processing.

Scenario 42 – Notebooks Chaining Through ADF

Explanation:

ADF executes multiple Databricks notebooks sequentially using *Execute Notebook Activity*.

Benefit: Creates logical layered data processing flows (Bronze→Silver→Gold).

Scenario 43 – ADF Calling Databricks REST API

Explanation:

Use Web activity to trigger Databricks jobs by API.

Benefit: Flexible for external triggers and CI/CD integration.

Scenario 44 – Parameter File from Blob

Explanation:

Databricks reads JSON config from Blob:

```
config=json.load(open("/dbfs/mnt/config/pipeline.json"))
```

Benefit: Config-driven ETL design.

Scenario 45 – Data Profiling Using Databricks Notebook

Explanation:

Generate summary stats like count, distinct, nulls.

Benefit: Helps data validation before load to warehouse.

Scenario 46 – Storing ADF Logs in Databricks

Explanation:

ADF sends run details to ADLS → Databricks aggregates for monitoring dashboard.

Benefit: Centralized pipeline analytics.

Scenario 47 – ADF and Databricks for Streaming Data

Explanation:

ADF pulls batch files while Databricks handles Spark Structured Streaming for near real-time processing.

Benefit: Unified batch + stream architecture.

Scenario 48 – Handling Schema Evolution in Delta Lake

Explanation:

Enable `spark.databricks.delta.schema.autoMerge.enabled=true` for new columns.

Benefit: Automatic schema adaptation without pipeline failures.

Scenario 49 – Databricks Job Failure Alert in ADF

Explanation:

Capture Activity Run Output status → If failed, trigger Logic App email/SMS.

Benefit: Immediate notification for support teams.

Scenario 50 – ADF + Databricks End-to-End ETL Example

Flow:

Blob → ADF copy to ADLS → Databricks transform → ADF load to Synapse.

Benefit: Complete enterprise-grade modern data pipeline.

Tips for Scenarios 31–50

- Stress how ADF and Databricks complement each other.
- Mention **Delta Lake**, **parameterization**, and **cost optimization** in answers.

- Be ready to draw flow diagrams if asked in interview.
-

Section 3 – ADF Real-Time & Error-Handling Scenarios (51–75)

Scenario 51 – Real-Time Data Ingestion Using Event Trigger

Description:

Automatically start pipelines when a new file lands.

Explanation:

Use *Event Grid trigger* on Blob container. It listens for file creation events and triggers ADF.

Benefit: Real-time, serverless ingestion—no need for manual scheduling.

Scenario 52 – Handling Late Arriving Files

Explanation:

Use a Wait + Get Metadata activity loop that checks file existence.

If not found, waits 10 mins and retries.

Benefit: Avoids pipeline failure due to delayed files.

Scenario 53 – Retry Mechanism on Failure

Explanation:

Configure activity-level *Retry count* and *Interval in seconds*.

Example: Retry Copy Activity 3 times if transient error occurs.

Benefit: Improves resiliency.

Scenario 54 – Error Logging to SQL Table

Explanation:

On Failure path → Stored Procedure activity → Log pipeline name, activity, and error message.

Benefit: Centralized error repository for analysis.

Scenario 55 – Skipping Failed Files

Explanation:

Loop through multiple files; use *If Condition* to skip file if Copy fails but continue pipeline.

Benefit: Avoids entire job failure due to one corrupt file.

Scenario 56 – Timeout Handling

Explanation:

Set *Timeout* property for long-running activities like Copy/Databricks.
If exceeded, pipeline fails gracefully.

Benefit: Controls resource utilization.

Scenario 57 – Validation Activity Before Load

Explanation:

Use Data Flow or Databricks to validate schema, nulls, duplicates before Copy.

Benefit: Prevents loading bad data into production systems.

Scenario 58 – Capturing Row Counts

Explanation:

Store Source Count and Target Count in pipeline variables or SQL audit table.
Compare counts post-load for data completeness.

Benefit: Ensures integrity of transferred data.

Scenario 59 – Parameter File for Dynamic Config

Explanation:

Read configuration JSON from Blob or SQL, then dynamically pass values to Linked Services, Datasets, and Pipelines.

Benefit: Single source of truth for environments.

Scenario 60 – Email Notification for Errors

Explanation:

Use Web Activity to trigger Logic App → Send email with Run ID, Activity Name, and Error Message.

Benefit: Real-time alerting to support teams.

Scenario 61 – Logging Success Runs

Explanation:

Add Stored Procedure activity at pipeline end → Log execution status as *Success*.

Benefit: End-to-end monitoring of all pipelines.

Scenario 62 – Detecting Missing Files

Explanation:

Use Lookup activity with expected filenames from SQL config.

Compare against *Get Metadata* results from Blob.

Benefit: Identifies missing feeds proactively.

Scenario 63 – Dynamic Retry with Web Activity

Explanation:

Use a custom Azure Function triggered from ADF to handle variable retry logic based on error code.

Benefit: Smart error recovery.

Scenario 64 – Capturing Pipeline Duration

Explanation:

Use system variables `@pipeline().RunStart` and `@utcNow()` to calculate duration.

Log it in SQL table.

Benefit: Helps performance analysis.

Scenario 65 – Archiving Processed Files

Explanation:

After successful load, use Copy Activity to move source file to `/archive/yyyy/MM/dd/`.

Then delete original.

Benefit: Maintains clean folder structure.

Scenario 66 – Dynamic Retry for Databricks Jobs

Explanation:

If Databricks notebook fails, catch error output → Retry after a delay → Stop after 3 attempts.

Benefit: Robust fault-tolerance for transient cluster issues.

Scenario 67 – Real-Time CDC using Azure SQL and ADF

Explanation:

Enable Change Data Capture (CDC) on Azure SQL → Use ADF to fetch incremental data using timestamp or version columns.

Benefit: Near real-time incremental loads.

Scenario 68 – Load Balancing Large File Transfers

Explanation:

Split large file into multiple chunks using *Copy parallelism* property.

Benefit: Faster throughput and efficient resource use.

Scenario 69 – Using Stored Procedure for Business Validation

Explanation:

ADF executes a stored procedure to check data rules before load.

If invalid → return error → pipeline stops.

Benefit: Integrates ADF with business validation logic.

Scenario 70 – Retry at Activity Group Level

Explanation:

Wrap multiple activities in an *Execute Pipeline Activity*.

Configure retry on this parent level.

Benefit: Simplifies error handling for grouped logic.

Scenario 71 – Logging ADF Run Details in Blob

Explanation:

Use Web Activity → Azure Function → Write JSON logs to Blob.

Benefit: Cloud-native logging without DB dependency.

Scenario 72 – Handling Null Columns Dynamically

Explanation:

Use *Derived Column* transformation in Data Flow to replace nulls with defaults.

Example: `iif(isNull(price), 0, price)`

Benefit: Prevents downstream load issues.

Scenario 73 – Debugging Copy Activity Performance

Explanation:

Enable *Logging Level = Verbose* to analyze throughput and bottlenecks.

Benefit: Optimizes pipeline design for better speed.

Scenario 74 – Using Global Parameters for Deployment

Explanation:

Define environment-wide parameters (e.g., region, resource group) at factory level.

Benefit: Simplifies CI/CD and environment-specific deployments.

Scenario 75 – Handling Duplicate Records in Load

Explanation:

Use Data Flow *Aggregate* or Databricks `dropDuplicates()` before loading to target.

Benefit: Ensures clean and unique data in destination.

Tips for Section 3:

- Focus on **real-time triggers, retries, and logging mechanisms**.
 - Mention **row count audits** and **CDC for incremental loads**.
 - Interviewers love hearing about **self-healing pipelines** and **metadata-driven validation**.
-

Section 4 – Advanced Enterprise & CI/CD Scenarios (76–100)

Scenario 76 – ADF CI/CD with Azure DevOps

Description:

Automate deployment between Dev, UAT, and Prod.

Explanation:

Use ADF's ARM Template export → Publish to DevOps repo → Create Release Pipeline to deploy via ARM Template.

Benefit: Promotes consistency and automation.

Scenario 77 – Parameterizing Environment Variables

Explanation:

Use global parameters like `EnvironmentName` and conditional expressions:

```
@if>equals(pipeline().globalParameters.Environment,'DEV'), 'DevStorage', 'ProdStorage')
```

Benefit: Single pipeline usable across all environments.

Scenario 78 – Blue-Green Deployment for ADF

Explanation:

Maintain two factories (ADF-Blue, ADF-Green). Deploy new changes to Green; switch routing post-

validation.

Benefit: Zero-downtime deployments.

Scenario 79 – ADF Integration with Git Repository

Explanation:

Connect ADF to GitHub/DevOps Git → Enable source control for pipelines and datasets.

Benefit: Versioning, rollback, and collaborative development.

Scenario 80 – Metadata-Driven Pipeline Framework

Explanation:

Create master pipeline → Lookup config table (source, target, load type) → ForEach executes Copy/Transform dynamically.

Benefit: Scalable and reusable data framework.

Scenario 81 – Dynamic Data Flow Creation

Explanation:

Use parameters in Data Flow for Source and Sink datasets.

Benefit: Single flow handles multiple data sources.

Scenario 82 – ADF Integration with Synapse Dedicated Pool

Explanation:

ADF loads processed data into Synapse tables via *Copy Activity* with PolyBase option.

Benefit: Fast bulk loading of large volumes.

Scenario 83 – ADF Integration with Azure Data Lake Gen2 Hierarchy

Explanation:

ADF dynamically writes data into hierarchical folders (year/month/day).

Benefit: Partitioned structure aids downstream analytics.

Scenario 84 – Optimizing ADF Copy Performance

Explanation:

Use:

- *Staging* with Blob for SQL to Synapse

- *Parallel Copy* for large tables
 - *PolyBase* for Synapse loading
- Benefit:** Achieves up to 10x faster loads.
-

Scenario 85 – Custom Logging Using Azure Function

Explanation:

ADF calls an Azure Function to log pipeline metadata (start, end, status).

Benefit: Unified logging across ADF and external processes.

Scenario 86 – Integrating ADF with Power Automate

Explanation:

Use Web Activity to call Power Automate flow for approval or manual validation step.

Benefit: Human-in-the-loop data validation.

Scenario 87 – Orchestrating ML Model Execution

Explanation:

ADF triggers Azure Machine Learning Pipeline or Databricks notebook to run models post data prep.

Benefit: Automates ML lifecycle integration.

Scenario 88 – Controlled Backfill Load

Explanation:

For historical data load, use pipeline parameters for start/end dates.

Loop through months and process each batch sequentially.

Benefit: Efficient backfill without overloading resources.

Scenario 89 – ADF Integration with REST API

Explanation:

Use REST Linked Service + Copy Activity to extract data from APIs like Salesforce.

Benefit: Simplifies data ingestion from SaaS sources.

Scenario 90 – Pagination in API Calls

Explanation:

Use dynamic URL:

```
@concat('https://api/data?page=', item())
```

Loop through pages using *Until Activity*.

Benefit: Handles large API data extraction.

Scenario 91 – ADF Trigger Dependency

Explanation:

Configure *Trigger dependency* to run Pipeline-B only after successful completion of Pipeline-A.

Benefit: Manages sequence without manual scheduling.

Scenario 92 – Using Managed Identity for Authentication

Explanation:

Enable Managed Identity → Assign Storage Blob/Data Factory Contributor role.

Benefit: Secure, password-less authentication.

Scenario 93 – Parallel Execution of Pipelines

Explanation:

Run multiple data loads in parallel using *ForEach concurrency*.

Benefit: Reduces total runtime for multi-source loads.

Scenario 94 – Conditional Execution of Pipelines

Explanation:

Use *If Condition Activity*:

```
@if(equals(variables('LoadType'), 'Full'), true, false)
```

Execute based on condition.

Benefit: Flexible logic control.

Scenario 95 – Data Drift Handling in Data Flow

Explanation:

Enable *Allow Schema Drift* → ADF dynamically adapts to schema changes.

Benefit: No need to redesign pipelines for minor schema updates.

Scenario 96 – Hierarchical JSON Flattening

Explanation:

Use Data Flow's *Flatten* transformation for nested JSONs.

Example: API returning nested orders and items.

Benefit: Simplifies complex JSON ingestion.

Scenario 97 – ADF Integration with Logic Apps for SLA Monitoring

Explanation:

ADF triggers Logic App if pipeline exceeds expected duration.

Benefit: SLA alerting for delayed jobs.

Scenario 98 – End-to-End Data Lineage Tracking

Explanation:

Log source, transformations, and destination details in metadata tables.

Benefit: Enhances transparency and auditability.

Scenario 99 – Pipeline Run Analysis via REST API

Explanation:

Use ADF Management REST API to fetch pipeline run history for dashboarding.

Benefit: Monitoring and analytics integration.

Scenario 100 – ADF Orchestration for Enterprise Data Lakehouse

Flow:

ADF orchestrates → Databricks processes data → Delta Lake stores results → Power BI visualizes → Alerts via Logic App.

Benefit: Full modern data stack automation.

Tips for Section 4:

- Emphasize **DevOps CI/CD, parameterization, metadata-driven design, and secure identity-based access**.
 - These topics are favorites in **senior-level interviews** for ADF & Azure Data Engineer roles.
-

Thank You...!!